

Benchmarking

Goals and Strategies

Di Jin

Basics

Use the most realistic benchmark you can get. It can (and should) also impact your system design.

- Domain-specific advantages in real-world examples.
 - Good unit tests
 - Tracing infrastructure
- Accurate representation of the problems
 - Scale
 - Access patterns
 - Bottlenecks (e.g., IO, network)
- Acquisition
 - Research papers (selected benchmarks)
 - Popular open-source projects
 - Research papers (artifact)

Be Complete and Be Frank

Ultimate goal: allowing readers to extrapolate what will happen with their workload.

- Show good things
 - Most ideal scenario/usecase
 - Best workload
- Show the bad things
 - When does the scalability stop
 - What workload will be negatively impacted
- Show the performance breakdown
 - Different component's overhead scales differently
 - Future research can aim to improve specific part
 - Readers can confirm their understanding

What to Benchmark?

Extraordinary claims require extraordinary evidence.

- Something extremely surprising needs a lot of evidence
 - e.g., my system uses a single core but it's faster than a multi-threaded solution
 - Is there a breakdown?
 - How does it work on different types of workload?
 - What's the CPU utilization?
 - What are the potential limits?
 - ...
- Something apparently true may require no benchmarking
 - e.g., my system adds metadata to pip packages for supply-chain security, it incurs no runtime memory overhead



Infrastructure

- Consider use managed server instances (e.g., CloudLab <https://www.cloudlab.us/>, AWS)
- Daemonize the benchmarking process (e.g., screen, tmux, nohup)
- Repeatable
 - Setup correct environment
 - Avoid stale states
 - *Seriously* consider things like Docker

The moment all benchmarks run without any issue is probably the last time you run them. Anticipate problems, and be prepared to solve them.

- Granularity: don't put all tests into one scripts
- Fast iteration
 - Have small scale benchmarks to help catch problems early (run in seconds)
 - Have big scale benchmarks for main result (run this remotely, maybe hours or days)

“Watch it fail first”

- Be very very paranoid
- Verify whether things work as expected
 - Rough performance numbers
 - Deliberately triggered failures
 - Workload with predictably different outcomes
 - Manual inspections (e.g., debuggers, logging)

- Push-button effort
 - The benchmarks will be run **a lot**.
 - If you are lucky, they will be run by other people as well.
- Separate data generation and presentation
 - Allows analysis and inspection without re-running the experiment
 - Allows experimentation with data presentation when writing the paper
- Only add abstractions as needed. Avoid over-design.

Some Points on Methodology

No Interference

Do not run other things on the test environment

- No personal laptop (GUI, notifications, daemons)
- No running web server
- No resource sharing (even when VM is appropriate)

Report average of **repeated** experiments, and report **standard deviation** or **confidence interval**

- Not acceptable to say 5 second improvement with 20 second variation
- ~~Pick a good one~~ (You would go to academic jail for this)
- ~~Repeat 64 times to bring P-value down~~
- Figure out why it's varying so much and fix it
 - Interference
 - Networking in non-networking benchmarks
 - CPU auto-scaling/temperature throttling

Get Some Results First

Get 3–5 benchmarks end-to-end first

- Avoid fundamental problems
 - Do not go three months into a project then realize that Python's GIL is a problem
- Prioritize important things: *"If a tree falls in a forest and no one is around to hear it, does it make a sound?"*
 - Solve problems impacting the benchmarks
 - If you believe something is important, find corresponding benchmark first



- Add additional tracing for to breakdown overhead
- Beware of the bottleneck of the system
 - e.g., don't hide CPU overhead under network bottleneck
- Differentiate latency and throughput

Further Readings

Systems Benchmarking Crimes by Gernot Heiser

<https://gernot-heiser.org/benchmarking-crimes.html>